

# Protected-Mode Interrupt Processing

Chapter 14  
S. Dandamudi

# Outline

- Introduction
- Taxonomy of interrupts
- Interrupt processing
- Exceptions
- Software interrupts
- File I/O
  - File descriptor
  - File pointer
  - File system calls
- Illustrative examples
  - Write a character to display
  - Read a string from the keyboard
  - File copy
- Hardware interrupts

# Introduction

- Interrupts alter a program's flow of control
  - Behavior is similar to a procedure call
    - Some significant differences between the two
- Interrupt causes transfer of control to an *interrupt service routine* (ISR)
  - ISR is also called a *handler*
- When the ISR is completed, the original program resumes execution
- Interrupts provide an efficient way to handle unanticipated events

# Interrupts vs. Procedures

## Interrupts

- Initiated by both *software* and *hardware*
- Can handle *anticipated* and *unanticipated* internal as well as external events
- ISRs or interrupt handlers are memory resident
- Use numbers to identify an interrupt service
- (E)FLAGS register is saved automatically

2005

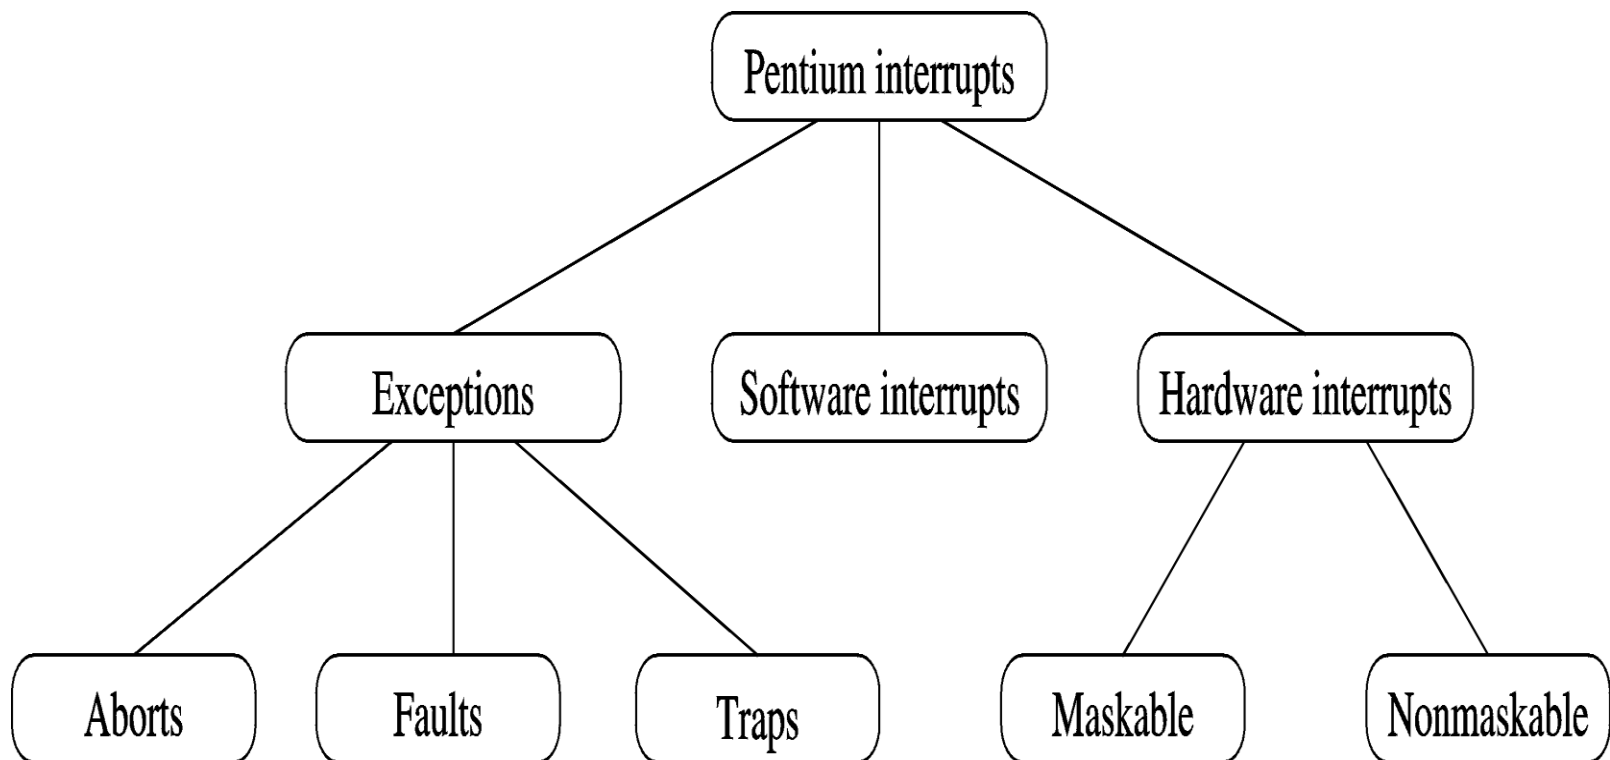
To be used with S.

Dandamudi, "Introduction to  
Assembly Language

## Procedures

- Can only be initiated by *software*
- Can handle *anticipated* events that are coded into the program
- Typically loaded along with the program
- Use meaningful names to indicate their function
- Do not save the (E)FLAGS register

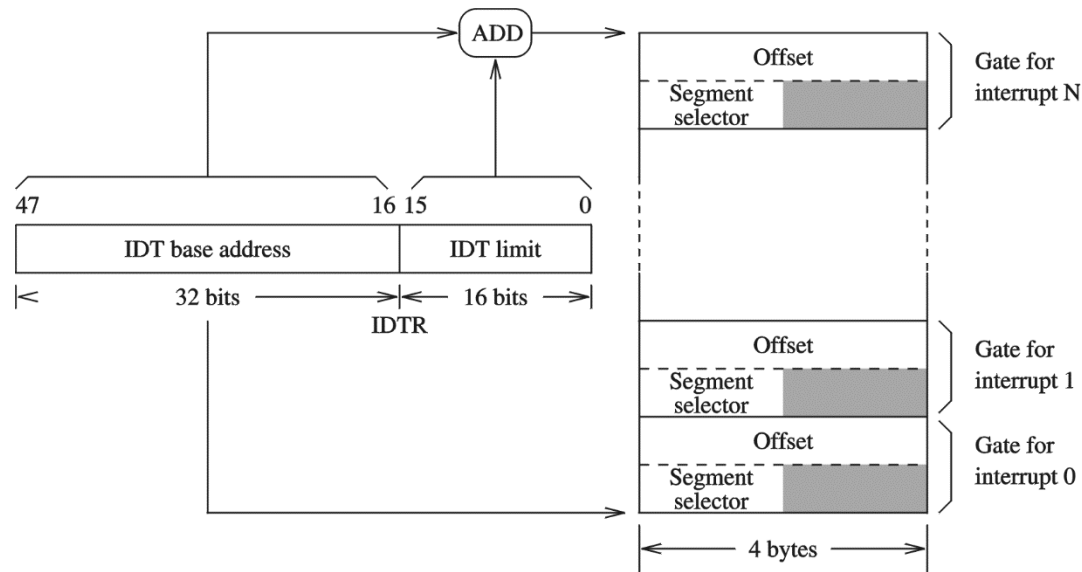
# A Taxonomy of Pentium Interrupts



# Protected Mode Interrupt Processing

- Up to 256 interrupts are supported (0 to 255)
  - Same number in both real and protected modes
  - Some significant differences between real and protected mode interrupt processing
- Interrupt number is used as an index into the *Interrupt Descriptor Table* (IDT)
  - This table stores the addresses of all ISRs
  - Each descriptor entry is 8 bytes long
    - Interrupt number is multiplied by 8 to get byte offset into IDT
  - IDT can be stored anywhere in memory
    - In contrast, real mode interrupt table has to start at address 0

# Protected Mode Interrupt Processing (cont'd)



## Organization of the IDT

2005

To be used with S.

Dandamudi, "Introduction to  
Assembly Language

S. Dandamudi

Chapter 14: Page 7

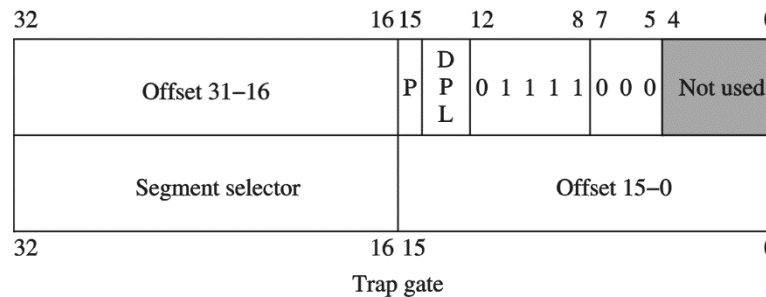
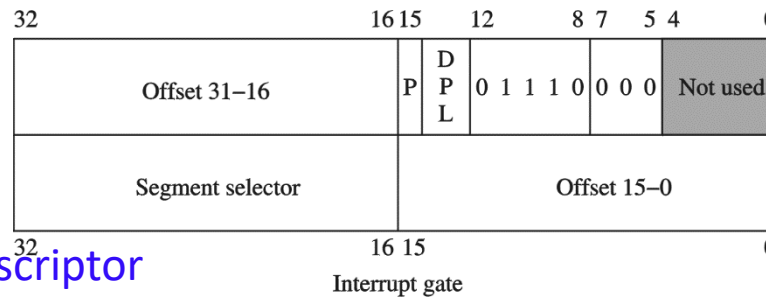
# Protected Mode Interrupt Processing (cont'd)

- Location of IDT is maintained by IDT register IDTR
- IDTR is a 48-bit register
  - 32 bits for IDT base address
  - 16 bits for IDT limit value
    - IDT requires only 2048 (11 bits)
    - A system may have smaller number of descriptors
      - Set the IDT limit to indicate the size in bytes
  - If a descriptor outside the limit is referenced
    - Processor enters shutdown mode
- Two special instructions to load (**lidt**) and store (**sidt**) IDT
  - Both take the address of a 6-byte memory as the operand

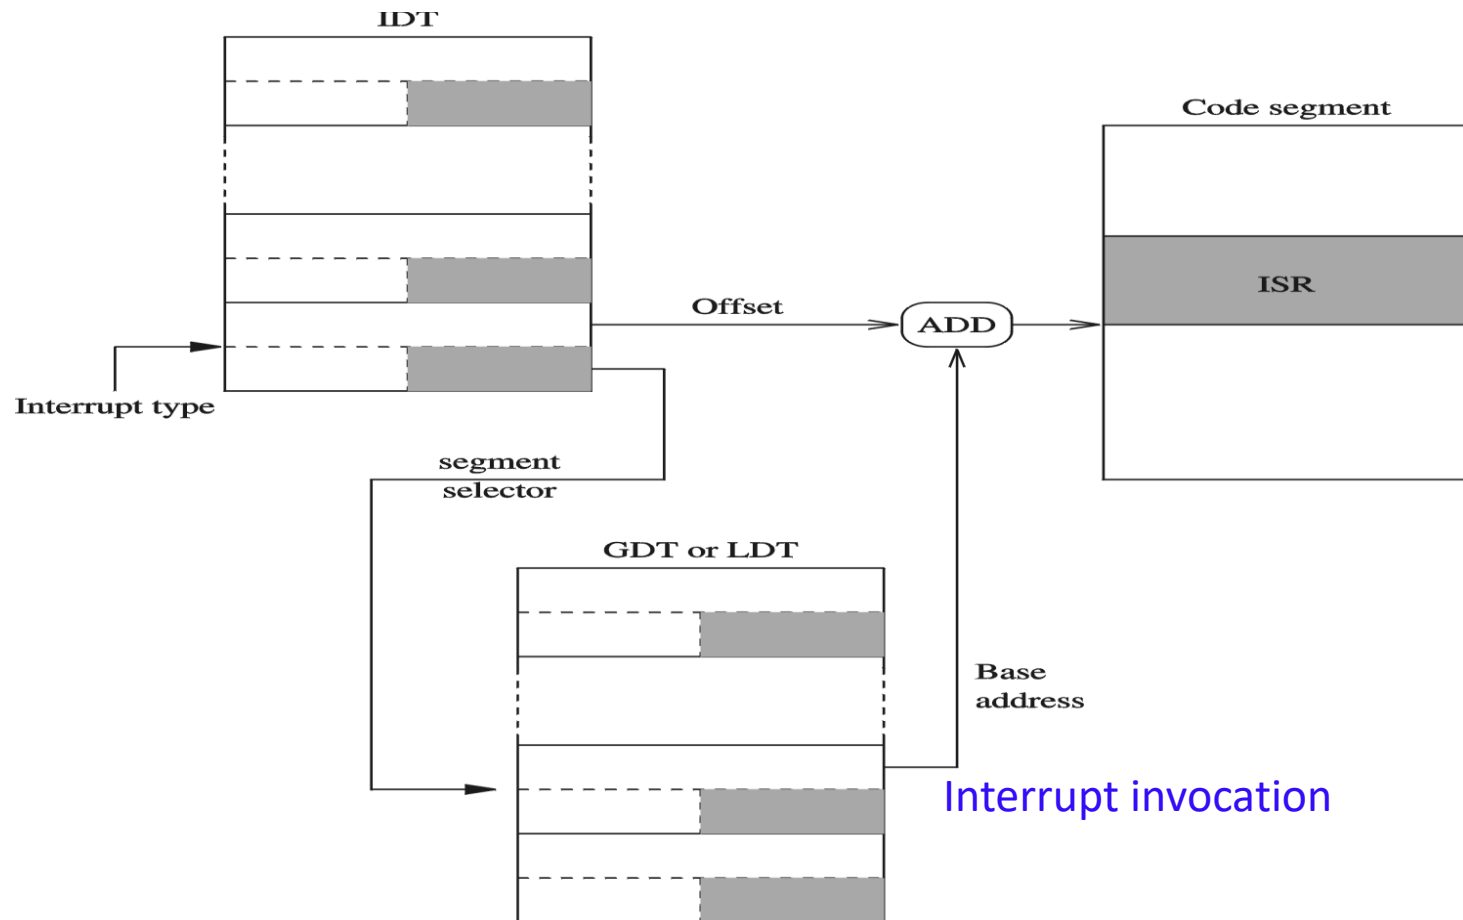


# Protected Mode Interrupt Processing (cont'd)

Interrupt descriptor



# Protected Mode Interrupt Processing (cont'd)



to be used with S.

# What Happens When An Interrupt Occurs?

- Push the EFLAGS register onto the stack
- Clear interrupt enable and trap flags
  - This disables further interrupts
  - Use `sti` to enable interrupts
- Push CS and EIP registers onto the stack
- Load CS with the 16-bit segment selector from the interrupt gate
- Load EIP with the 32-bit offset value from the interrupt gate

# Interrupt Enable Flag Instructions

- Interrupt enable flag controls whether the processor should be interrupted or not
- Clearing this flag disables all further interrupts until it is set
  - Use **c*l*i** (clear interrupt) instruction for this purpose
  - It is cleared as part interrupt processing
- Unless there is special reason to block further interrupts, enable interrupts in your ISR
  - Use **s*t*i** (set interrupt) instruction for this purpose

# Returning From An ISR

- As in procedures, the last instruction in an ISR should be **iret**
- The actions taken on **iret** are:
  - pop the 32-bit value on top of the stack into EIP register
  - pop the 16-bit value on top of the stack into CS register
  - pop the 32-bit value on top of the stack into the EFLAGS register
- As in procedures, make sure that your ISR does not leave any data on the stack
  - Match your push and pop operations within the ISR

# Exceptions

- Three types of exceptions
  - Depending on the way they are reported
  - Whether or not the interrupted instruction is restarted
    - Faults
    - Traps
    - Aborts
- Faults and traps are reported at instruction boundaries
- Aborts report severe errors
  - Hardware errors
  - Inconsistent values in system tables

# Faults and Traps

- Faults

- Instruction boundary before the instruction during which the exception was detected
- Restarts the instruction
- Divide error (detected during **div/ldiv** instruction)
- Segment-not-found fault

- Traps

- Instruction boundary immediately after the instruction during which the exception was detected
- No instruction restart
- Overflow exception (interrupt 4) is a trap
- User defined interrupts are also examples of traps

# Dedicated Interrupts

- Several Pentium predefined interrupts --- called dedicated interrupts
- These include the first five interrupts:

interrupt type	Purpose
0	Divide error
1	Single-step
2	Nonmaskable interrupt (MNI)
3	Breakpoint
4	Overflow



# Dedicated Interrupts (cont'd)

- Divide Error Interrupt
  - CPU generates a type 0 interrupt whenever the div/ldiv instructions result in a quotient that is larger than the destination specified
- Single-Step Interrupt
  - Useful in debugging
  - To single step, Trap Flag (TF) should be set
  - CPU automatically generates a type 1 interrupt after executing each instruction if TF is set
  - Type 1 ISR can be used to present the system state to the user

# Dedicated Interrupts (cont'd)

- Breakpoint Interrupt
  - Useful in debugging
  - CPU generates a **type 3** interrupt
  - Generated by executing a special single-byte version of **int 3** instruction (opcode CCH)
- Overflow Interrupt
  - Two ways of generating this type 4 interrupt
    - **int 4** (unconditionally generates a type 4 interrupt)
    - **into** (interrupt is generated only if the overflow flag is set)
  - We do not normally use **into** as we can use **jo/jno** conditional jumps to take care of overflow

# Software Interrupts

- Initiated by executing an interrupt instruction

**int      interrupt-type**

**interrupt-type** is an integer in the range 0 to 255

- Each interrupt type can be parameterized to provide several services.
- For example, Linux interrupt service **int    0x80** provides a large number of services (more than 180 system calls!)
  - EAX register is used to identify the required service under **int 0x80**

# File I/O

- Focus is on File I/O
  - Keyboard and display are treated as stream files
  - Three standard file streams are defined
    - Standard input (**stdin**)
      - Associated device: Keyboard
    - Standard output (**stdout**)
      - Associated device: Display
    - Standard error (**stderr**)
      - Associated device: Display

# File I/O (cont'd)

- File descriptor
  - Small integer acts as a file id
  - Use file descriptors to access open files
  - File descriptor is returned by the **open** and **create** systems calls
  - Don't have to open the three standard files
    - Lowest three integers are assigned to these files
      - **stdin** (0)
      - **stdout** (1)
      - **stderr** (2)

# File I/O (cont'd)

- File pointer
  - Associated with each open file
  - Specifies offset (in bytes) relative to the beginning of the file
    - Read and write operations use this location
  - When a file is opened, file pointer points to the first byte
    - Subsequent reads move it to facilitate sequential access
  - Direct access to a file
    - Can be provided by manipulating the file pointer

# File System Calls

- File create call

## System call 8 --- Create and open a file


Inputs: EAX = 8

EBX = file name

ECX = file permissions

Returns: EAX = file descriptor

Error: EAX = error code



2005

To be used with S.

Dandamudi, "Introduction to  
Assembly Language

S. Dandamudi

Chapter 14: Page 23

# File System Calls (cont'd)

- File open call

## System call 5 --- Open a file

Inputs: EAX = 5

EBX = file name

ECX = file access mode

EDX = file permissions

Returns: EAX = file descriptor

Error: EAX = error code



# File System Calls (cont'd)

- File read call

## System call 3 --- Read from a file

Inputs: EAX = 3

EBX = file descriptor

ECX = pointer to input buffer

EDX = buffer size

(max. # of bytes to read)

Returns: EAX = # of bytes read

Error: EAX = error code

# File System Calls (cont'd)

- File write call

## System call 4 --- Write to a file

Inputs: EAX = 4

EBX = file descriptor

ECX = pointer to output buffer

EDX = buffer size

(# of bytes to write)

Returns: EAX = # of bytes written

Error: EAX = error code

# File System Calls (cont'd)

- File close call

## System call 6 --- Close a file

Inputs: EAX = 6

EBX = file descriptor

Returns: EAX = ---

Error: EAX = error code

# File System Calls (cont'd)

- File seek call

**System call 19 --- lseek (updates file pointer)**

**Inputs:** EAX = 19

EBX = file descriptor

ECX = offset

EDX = whence

**Returns:** EAX = byte offset from the  
beginning of file

**Error:** EAX = error code

# File System Calls (cont'd)

- **whence** value

Reference position	whence value
Beginning of file	0
Current position	1
End of file	2

# Illustrative Examples

- Three examples
  - Write a character to display
    - **putch** procedure
  - Read a string from the keyboard
    - **getstr** procedure
  - File copy
    - **file\_copy.asm**

# Hardware Interrupts

- Software interrupts are synchronous events
  - Caused by executing the `int` instruction
- Hardware interrupts are of hardware origin and asynchronous in nature
  - Typically caused by applying an electrical signal to the processor chip
- Hardware interrupts can be
  - Maskable
  - Non-maskable
    - Causes a `type 2` interrupt

# How Are Hardware Interrupts Triggered?

- Non-maskable interrupt is triggered by applying an electrical signal to the MNI pin of processor
  - Processor always responds to this signal
  - Cannot be disabled under program control
- Maskable interrupt is triggered by applying an electrical signal to the INTR (INTerrupt Request) pin of Pentium
  - Processor recognizes this interrupt only if IF (interrupt enable flag) is set
  - Interrupts can be masked or disabled by clearing IF



# How Does the CPU Know the Interrupt Type?

- Interrupt invocation process is common to all interrupts
  - Whether originated in software or hardware
- For hardware interrupts, processor initiates an interrupt acknowledge sequence
  - processor sends out interrupt acknowledge (INTA) signal
  - In response, interrupting device places interrupt vector on the data bus
  - Processor uses this number to invoke the ISR that should service the device (as in software interrupts)

# How can More Than One Device Interrupt?

- Processor has only one INTR pin to receive interrupt signal
- Typical system has more than one device that can interrupt --- keyboard, hard disk, floppy, etc.
- Use a special chip to prioritize the interrupts and forward only one interrupt to the CPU
  - 8259 Programmable Interrupt Controller chip performs this function (more details in Chapter 15)